

DATA DAYS

# Semi-supervised learning with GANs

Olga Petrova  
Machine Learning Engineer @ Scaleway



# OUTLINE

## THEORY :

- Semi-Supervised Machine Learning
- Generative Adversarial Networks (GANs)
- Semi-supervised Learning with GANs

## PRACTICE :

- A TALE OF CATS & DOGS

# THEORY

## Semi-Supervised Machine Learning

# WHAT IS SEMI-SUPERVISED LEARNING?

UNSUPERVISED LEARNING (unlabeled training data): e.g. clustering, GANs

SUPERVISED LEARNING (labeled training data): e.g. regression, classification

SEMI-SUPERVISED LEARNING:

- a subset of the training data is labeled
- the rest is unlabeled
- supervised learning task (e.g. classification)

# WHAT IS SEMI-SUPERVISED LEARNING?

**UNSUPERVISED LEARNING** (unlabeled training data): e.g. clustering, GANs

**SUPERVISED LEARNING** (labeled training data): e.g. regression, classification

**SEMI-SUPERVISED LEARNING:**

- a subset of the training data is labeled
- the rest is unlabeled
- supervised learning task (e.g. classification)

# WHAT IS SEMI-SUPERVISED LEARNING?

**UNSUPERVISED LEARNING** (unlabeled training data): e.g. clustering, GANs

**SUPERVISED LEARNING** (labeled training data): e.g. regression, classification

**SEMI-SUPERVISED LEARNING:**

- a subset of the training data is labeled
- the rest is unlabeled
- supervised learning task (e.g. classification)

# WHAT IS SEMI-SUPERVISED LEARNING?

**UNSUPERVISED LEARNING** (unlabeled training data): e.g. clustering, GANs

**SUPERVISED LEARNING** (labeled training data): e.g. regression, classification

**SEMI-SUPERVISED LEARNING:**

- a subset of the training data is labeled
- the rest is unlabeled
- supervised learning task (e.g. classification)

# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: supervised machine learning
- Good news: today's supervised models are very powerful
- But: often comes at the expense of very large/deep networks
- Bad news: the larger the model, the more data it needs to train!
- Really bad news: labelling data is expensive
- Unlabeled data to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models

# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- Good news: today's supervised models are very powerful
- But: often comes at the expense of very large/deep networks
- Bad news: the larger the model, the more data it needs to train!
- Really bad news: labelling data is expensive
- Unlabeled data to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models

# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- 😺 Good news: today's supervised models are **very powerful**
- But: often comes at the expense of **very large/deep networks**
- Bad news: the larger the model, the more data it needs to train!
- Really bad news: **labelling data is expensive**
- Unlabeled data to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models

# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- 😺 Good news: today's supervised models are **very powerful**
- 😺 But: often comes at the expense of **very large/deep networks**
- Bad news: the larger the model, the more data it needs to train!
- Really bad news: labelling data is expensive
- Unlabeled data to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models

# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- 😺 Good news: today's supervised models are **very powerful**
- 😺 But: often comes at the expense of **very large/deep networks**
- 😺 Bad news: the larger the model, **the more data it needs** to train!
- Really bad news: labelling data is expensive
- Unlabeled data to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models

# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- 😺 Good news: today's supervised models are **very powerful**
- 😺 But: often comes at the expense of **very large/deep networks**
- 😺 Bad news: the larger the model, **the more data it needs** to train!
- 😺 Really bad news: **labelling data is expensive**
- Unlabeled data to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models

# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- 😺 Good news: today's supervised models are **very powerful**
- 😺 But: often comes at the expense of **very large/deep networks**
- 😺 Bad news: the larger the model, **the more data it needs** to train!
- 😺 Really bad news: **labelling data is expensive** especially in highly specialized domains e.g. medicine
- Unlabeled data to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models



# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- 😺 Good news: today's supervised models are **very powerful**
- 😺 But: often comes at the expense of **very large/deep networks**
- 😺 Bad news: the larger the model, **the more data it needs** to train!
- 😺 Really bad news: **labelling data is expensive**
- **Unlabeled data** to the rescue
- Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models

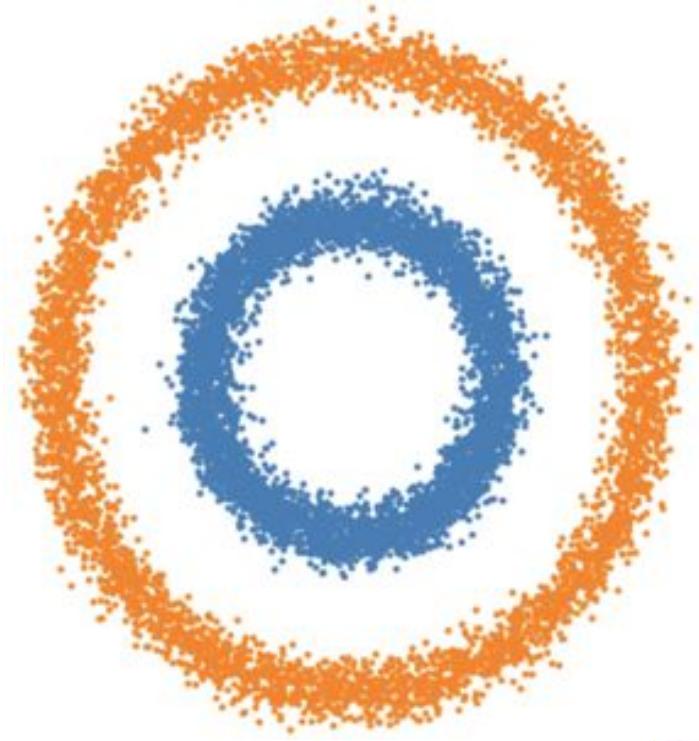


# WHY DO SEMI-SUPERVISED LEARNING?

- Most practical tasks: **supervised machine learning**
- 😺 Good news: today's supervised models are **very powerful**
- 😺 But: often comes at the expense of **very large/deep networks**
- 😺 Bad news: the larger the model, **the more data it needs** to train!
- 😺 Really bad news: **labelling data is expensive**
- **Unlabeled data** to the rescue
- **Semi-supervised learning: use unlabelled data to boost the performance of supervised ML models**

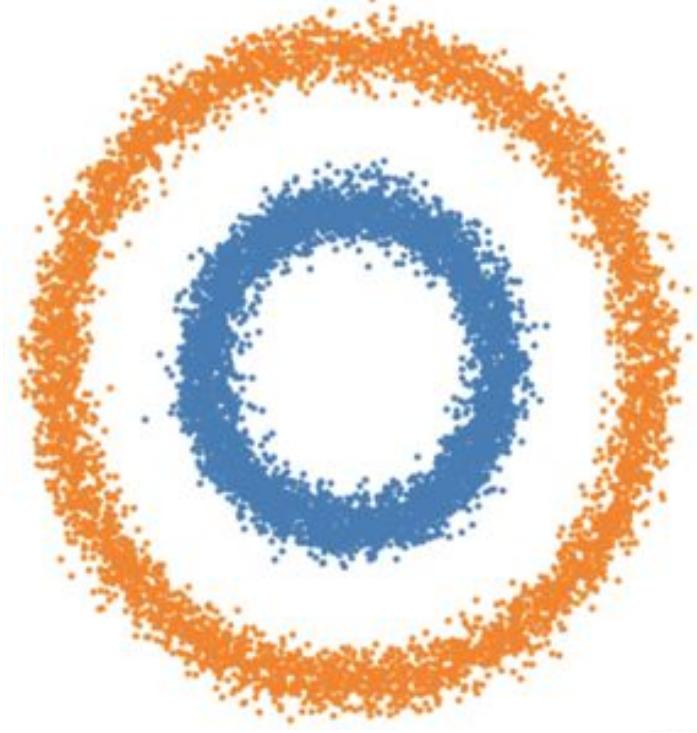


# HOW DOES SSL WORK?



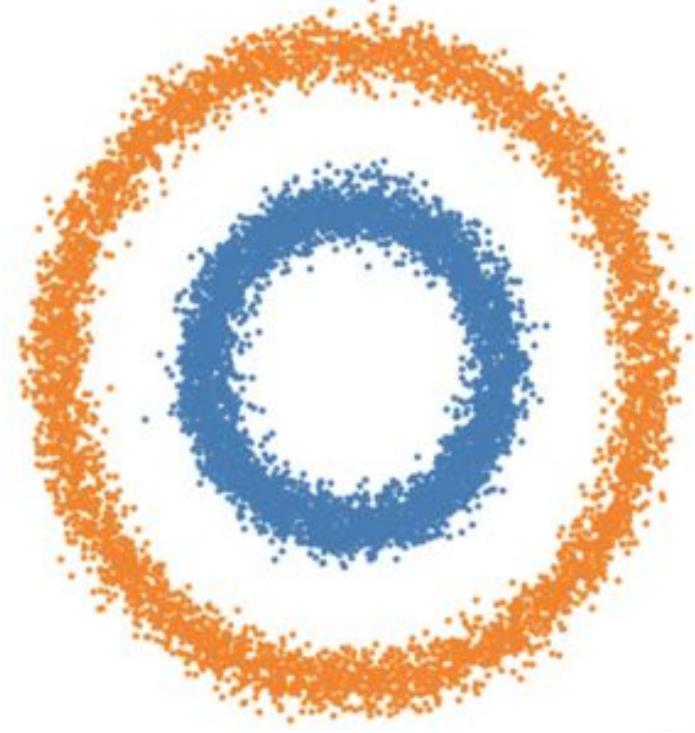
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: blues and oranges
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



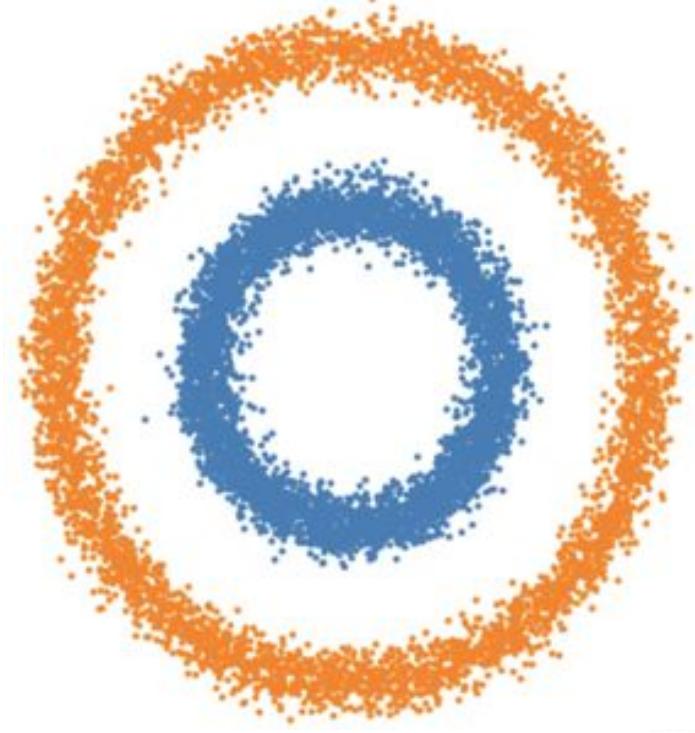
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



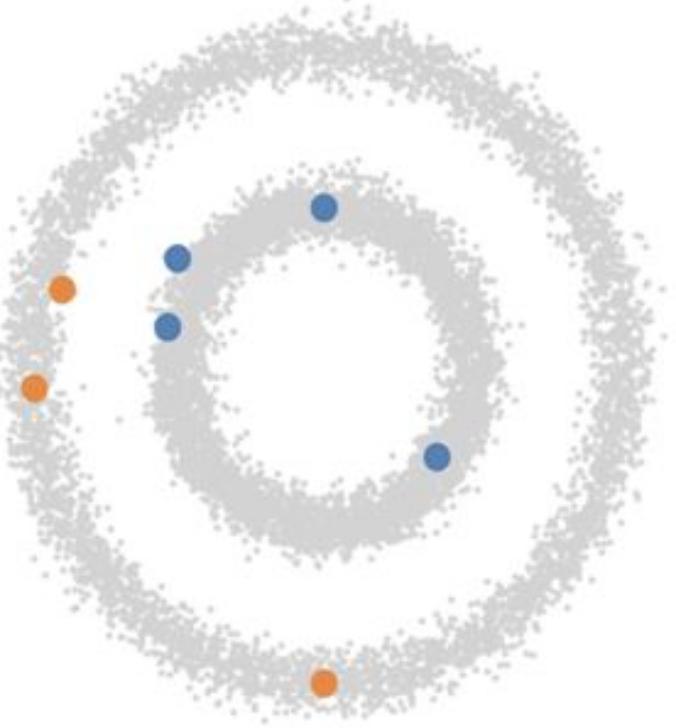
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



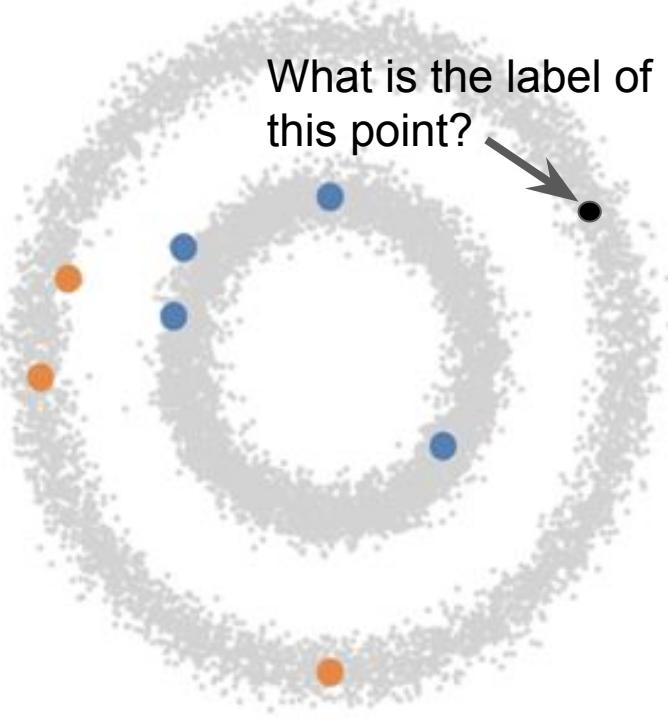
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



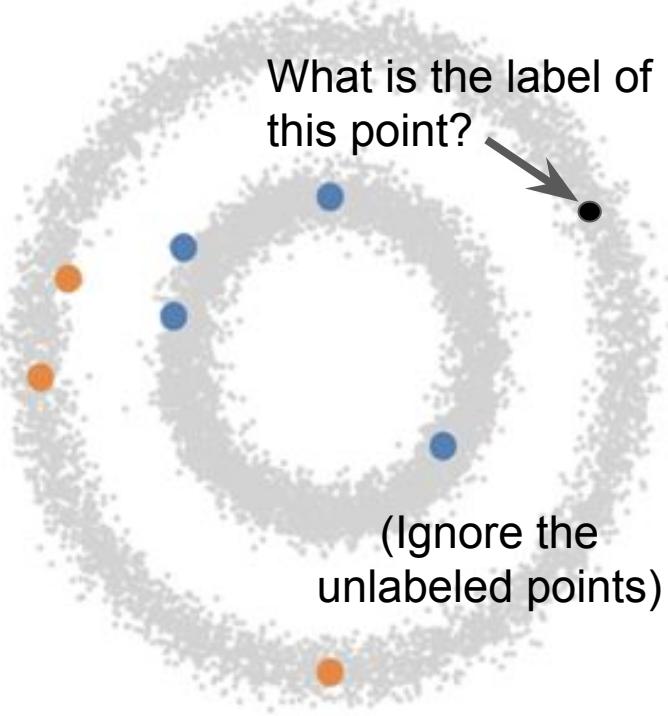
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



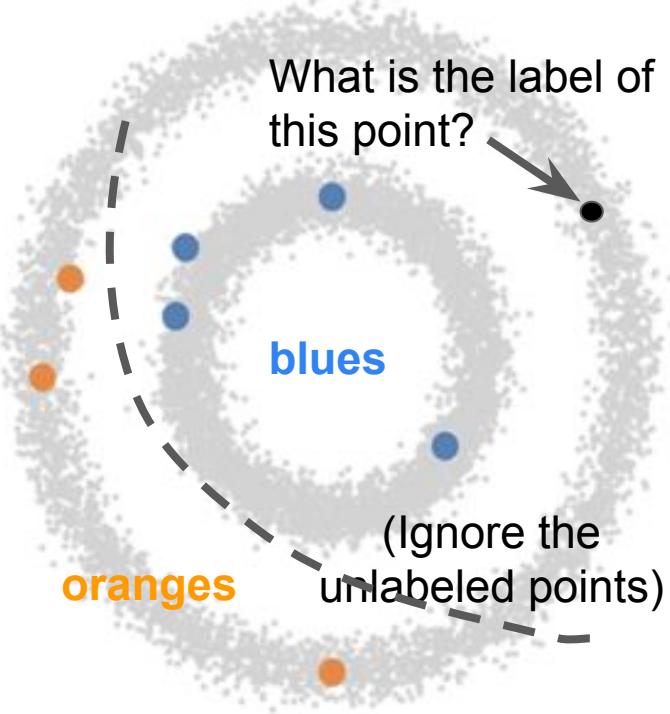
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



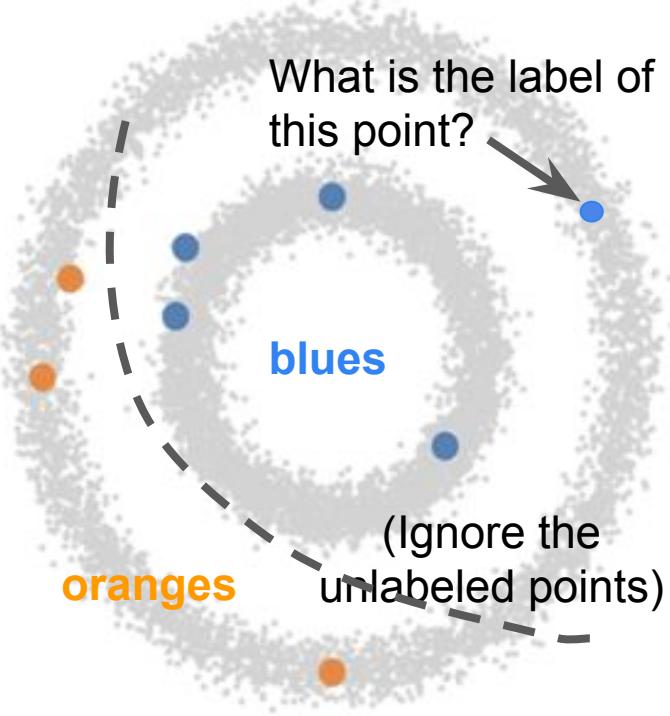
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



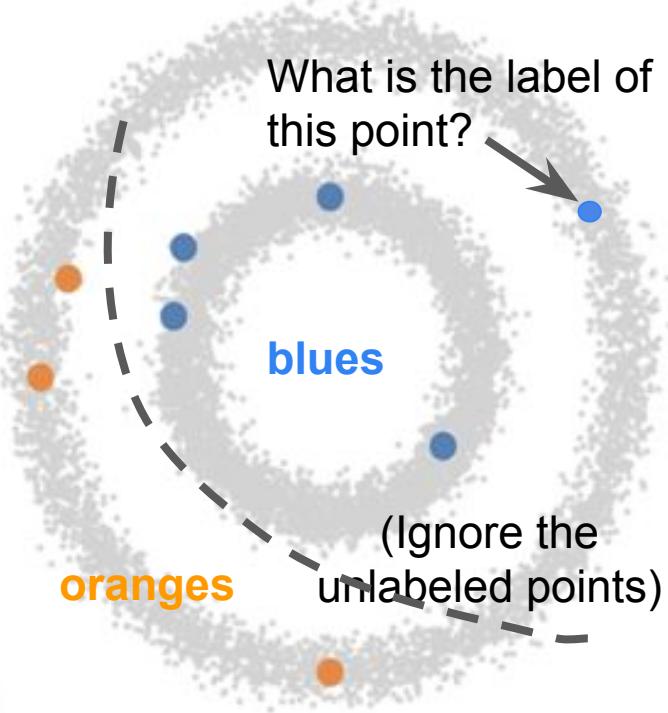
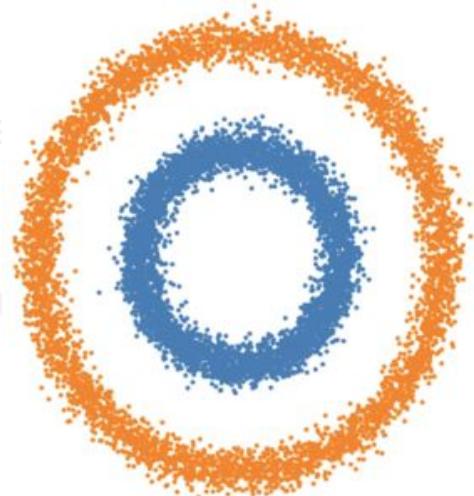
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



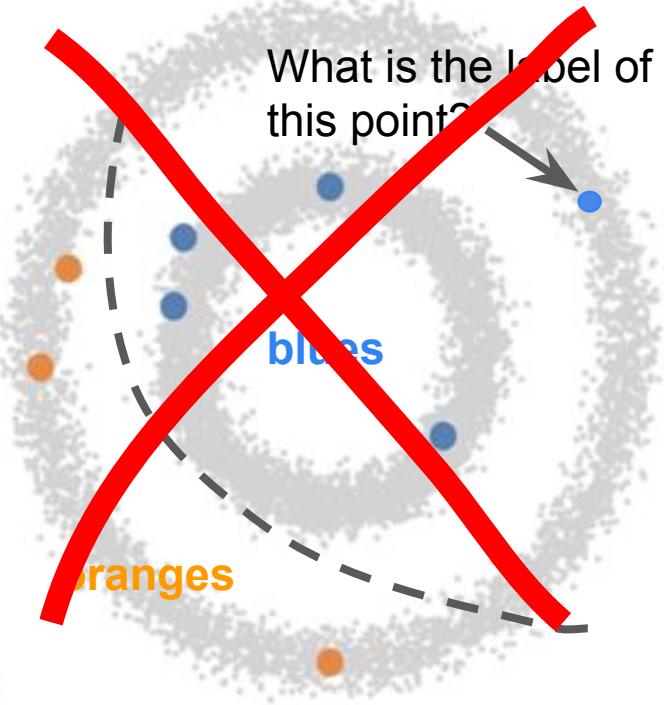
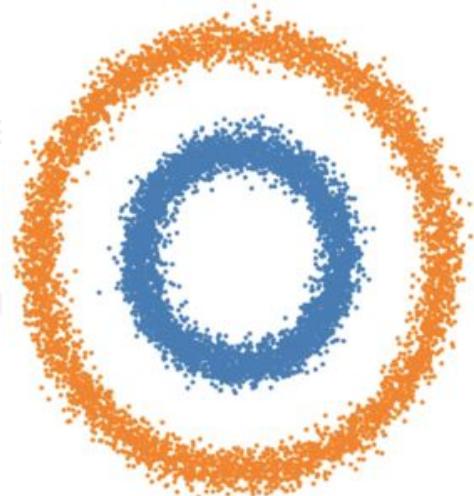
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is the data forms two rings
- With labeled points, we can assign labels to the rings



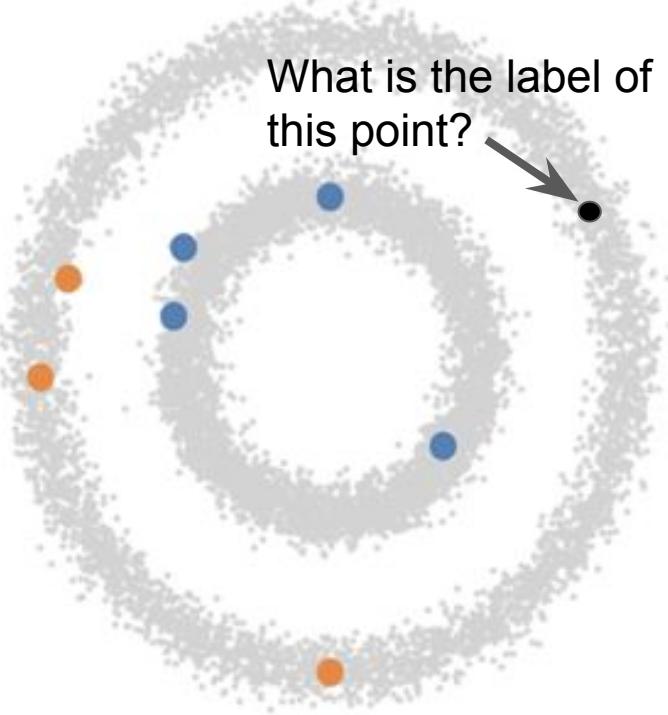
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is the data forms two rings
- With labeled points, we can assign labels to the rings



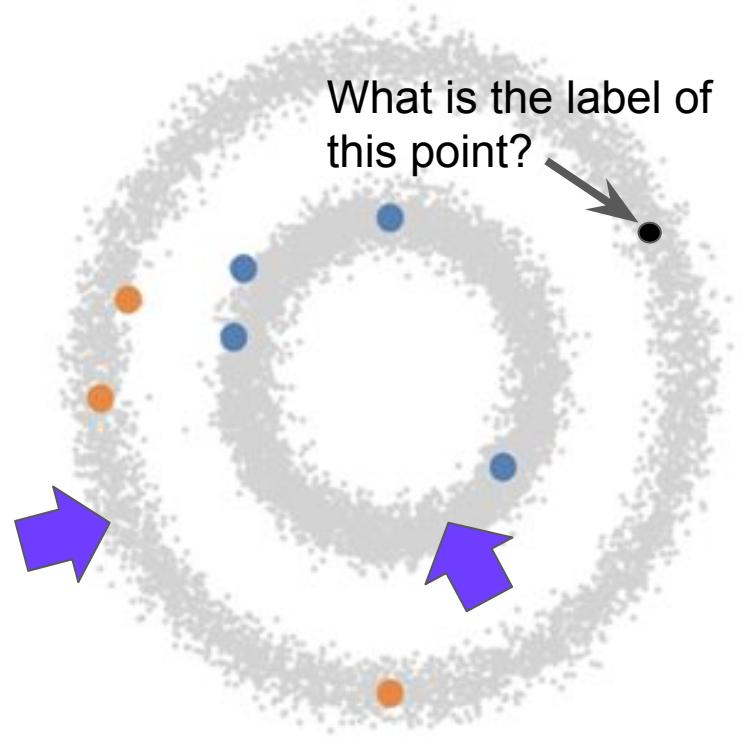
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



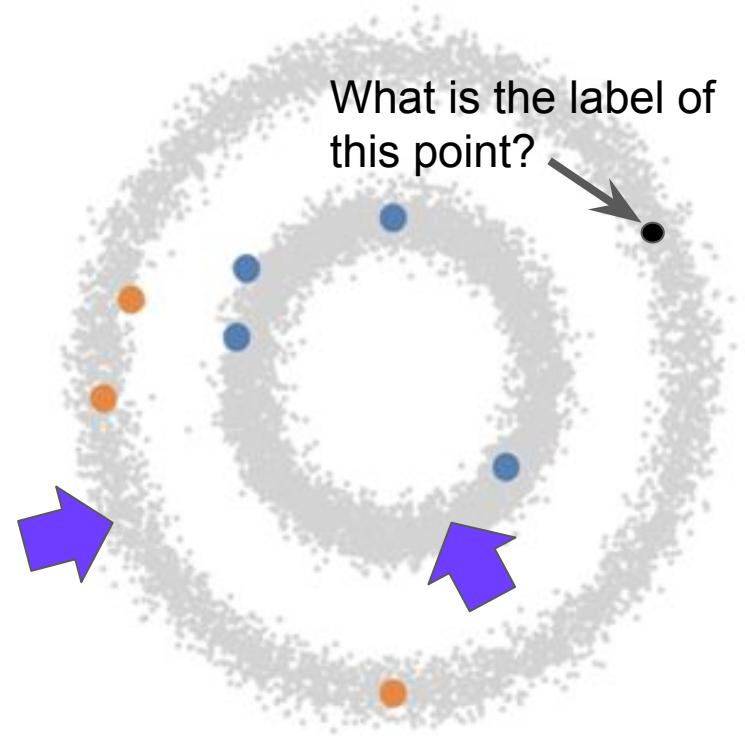
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



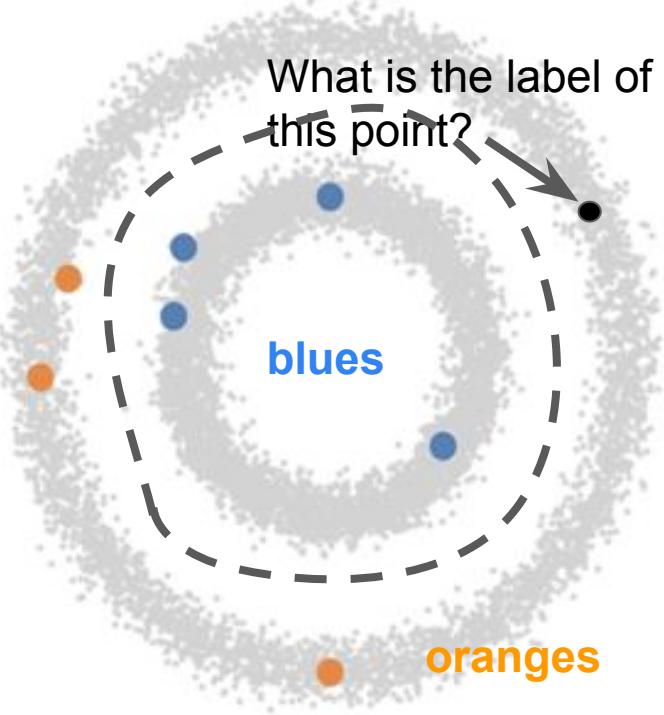
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



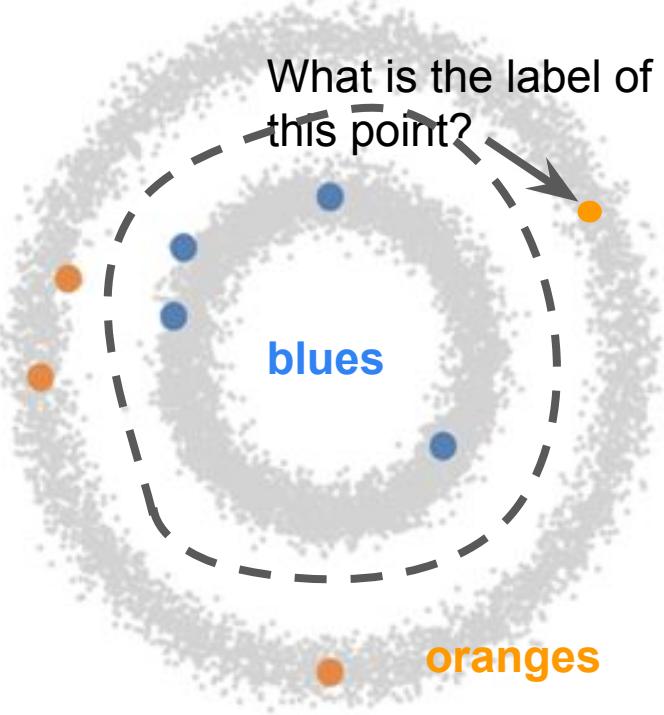
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



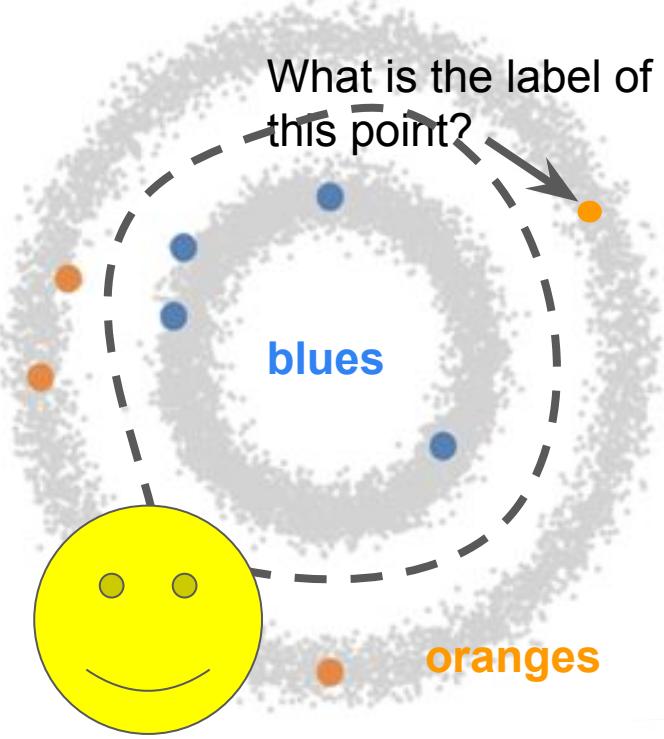
# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings



# HOW DOES SSL WORK?

- Consider a simple classification task
- Two classes: **blues** and **oranges**
- What if we only had a few training samples?
- With unlabeled points, it is easy to see that the data forms two rings
- With labeled points, we can assign class labels to the rings

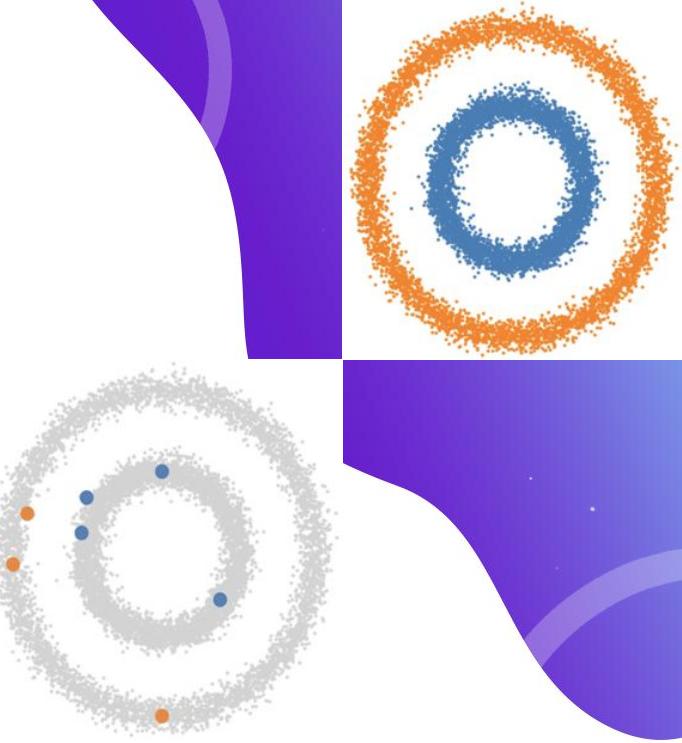


# HOW DOES SSL WORK?

## The Idea

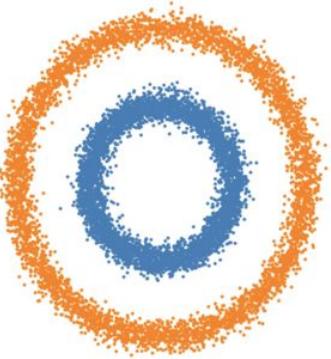
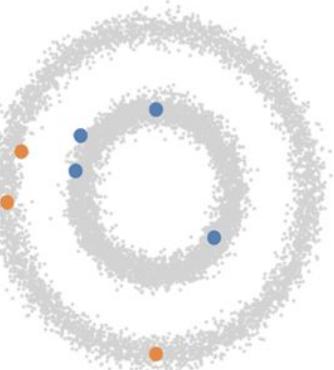
The unlabeled points give us information about the distribution of the data (e.g. through clustering or GANs)

The labeled points can then be used to assign class labels



# WHAT ARE SOME OF THE WAYS TO DO SSL?

- TODO



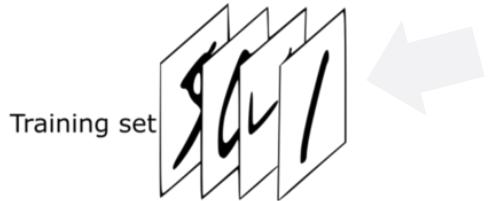
# THEORY

## Generative Adversarial Networks

# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Unsupervised learning: unlabelled training data (e.g., MNIST images)
- Goal: generate data from the same distribution (e.g. new images that look like MNIST)



[Image source](#)

# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Unsupervised learning: unlabelled training data (e.g., MNIST images)
- Goal: generate data from the same distribution (e.g. new images that look like MNIST)



[Image source](#)

# GANs: Generative Adversarial Networks

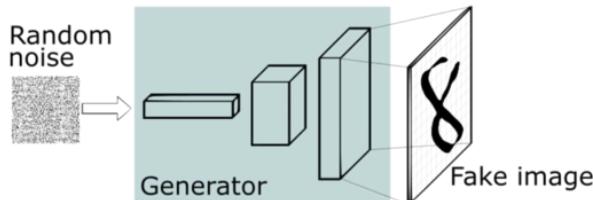
Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: Generator + Discriminator

# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise



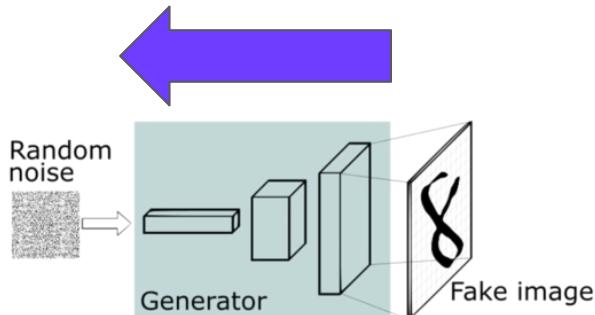
# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise

Most image classifiers: **convolutional network**  
Takes an image, reduces its size at each layer  
Image → a 1D array (vector)

**Deconvolutional network:** the opposite mechanism  
Start with a vector, increase the size at each layer  
Input vector → an Image



[Image source](#)

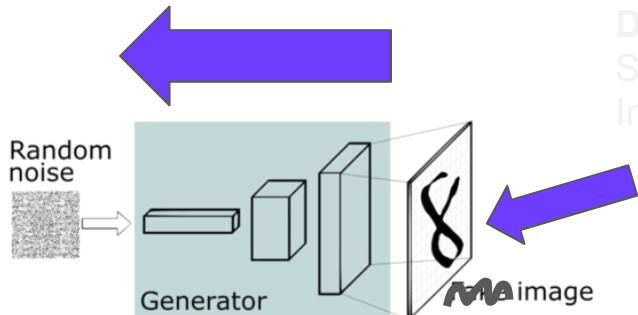
# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise

Most image classifiers: **convolutional network**  
Takes an image, reduces its size at each layer  
Image → a 1D array (vector)

Deconvolutional network: the opposite mechanism  
Start with a vector, increase the size at each layer  
Input vector → an Image



[Image source](#)

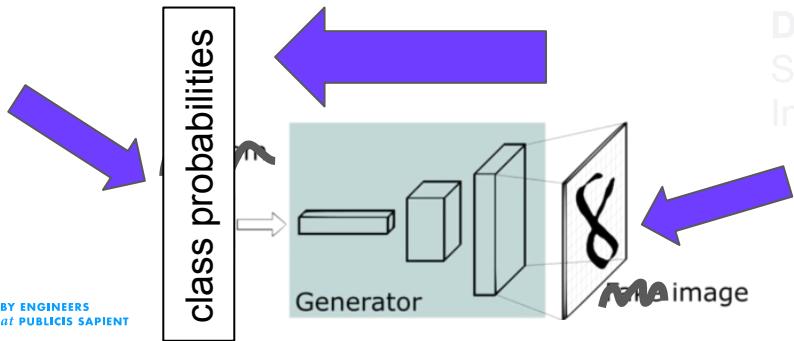
# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise

Most image classifiers: **convolutional network**  
Takes an image, reduces its size at each layer  
Image → a 1D array (vector)

Deconvolutional network: the opposite mechanism  
Start with a vector, increase the size at each layer  
Input vector → an Image



[Image source](#)

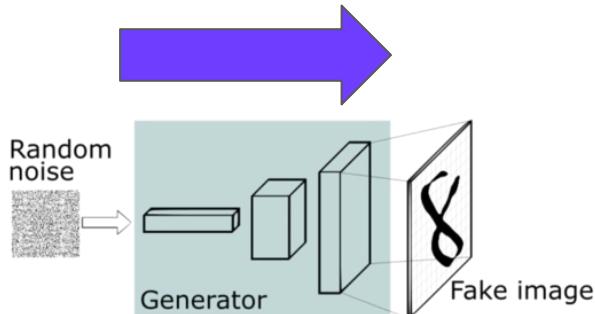
# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise

Most image classifiers: **convolutional network**  
Takes an image, reduces its size at each layer  
Image → a 1D array (vector)

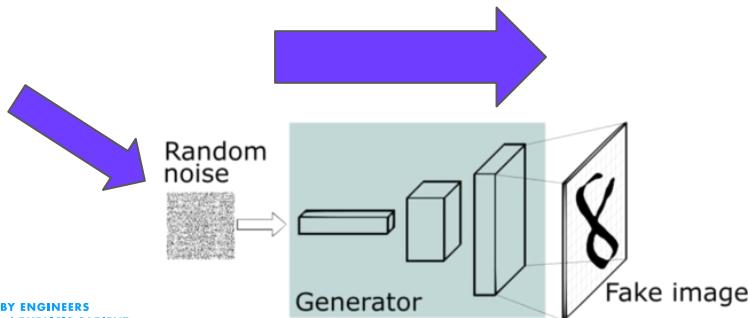
**Deconvolutional network:** the opposite mechanism  
Start with a vector, increase the size at each layer  
Input vector → an Image



# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise



Most image classifiers: **convolutional network**  
Takes an image, reduces its size at each layer  
Image → a 1D array (vector)

**Deconvolutional network:** the opposite mechanism  
Start with a vector, increase the size at each layer  
Input vector → an Image

[Image source](#)

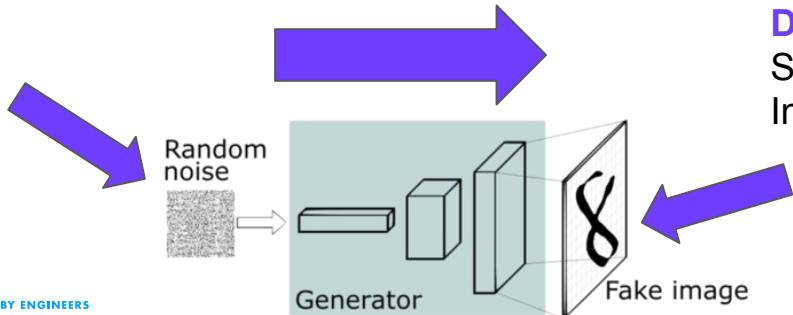
# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise

Most image classifiers: **convolutional network**  
Takes an image, reduces its size at each layer  
Image → a 1D array (vector)

**Deconvolutional network:** the opposite mechanism  
Start with a vector, increase the size at each layer  
Input vector → an Image

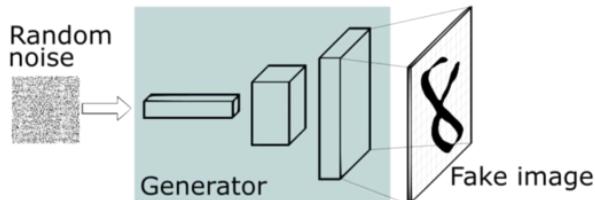


[Image source](#)

# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

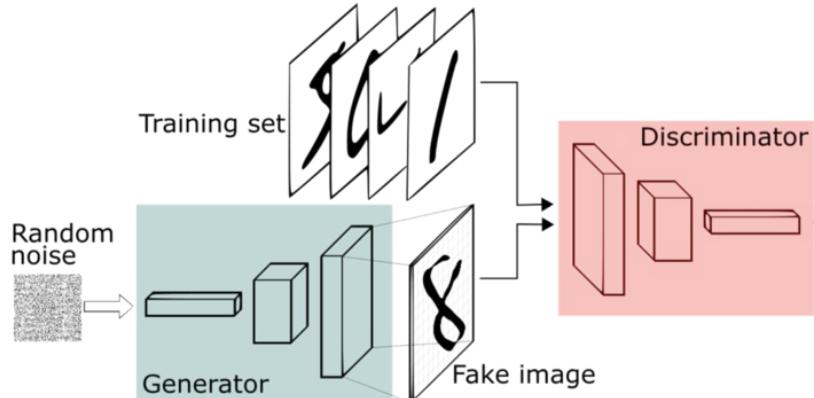
- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise



# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise
- **Discriminator**

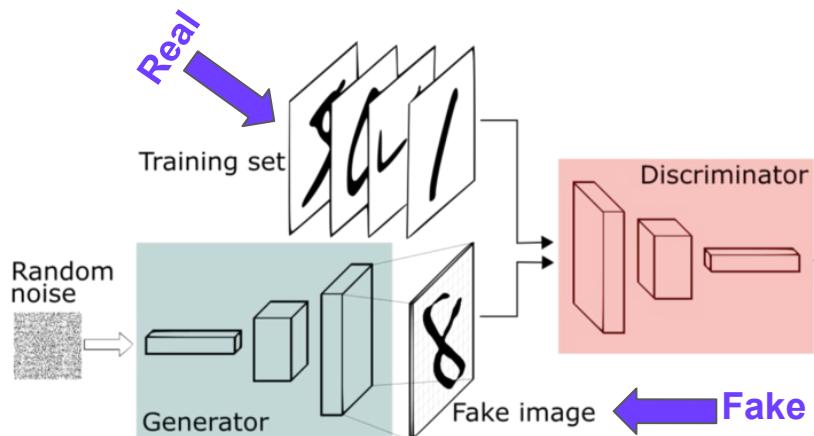


[Image source](#)

# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator + Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise
- **Discriminator**

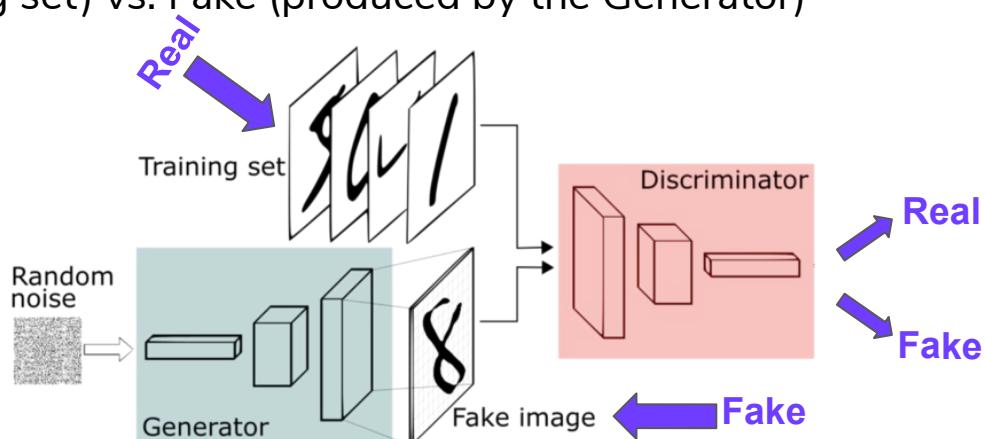


[Image source](#)

# GANs: Generative Adversarial Networks

Ref: [Ian Goodfellow et al., 2014](#)

- Two networks: **Generator** + **Discriminator**
- **Generator:** deconvolutional network  
generates an image from random input noise
- **Discriminator:** binary classifier trained to detect Real (from the training set) vs. Fake (produced by the Generator)



[Image source](#)

# GANs: Generative Adversarial Networks

## Training the Discriminator

- Straightforward binary classification (Real vs. Fake)

## Training the Generator

- Generator produces an image → pass it through the Discriminator
- If Discriminator classifies it as Fake → update the Generator to produce better images

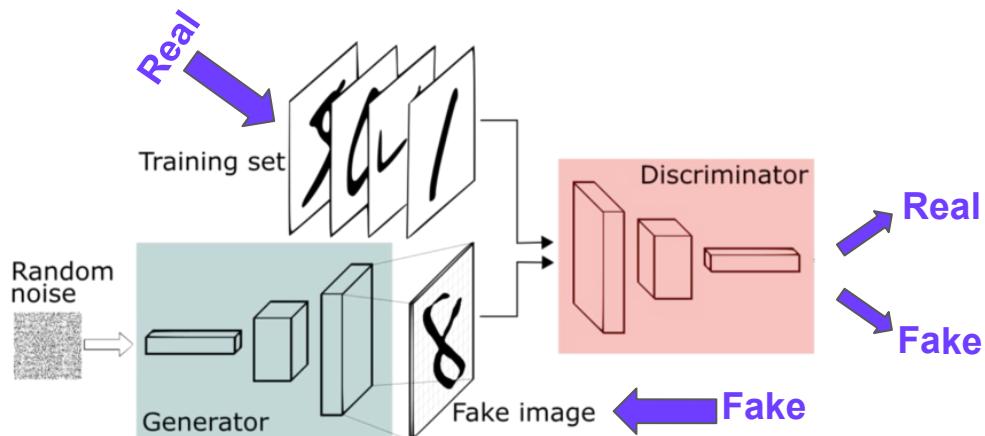


Image source

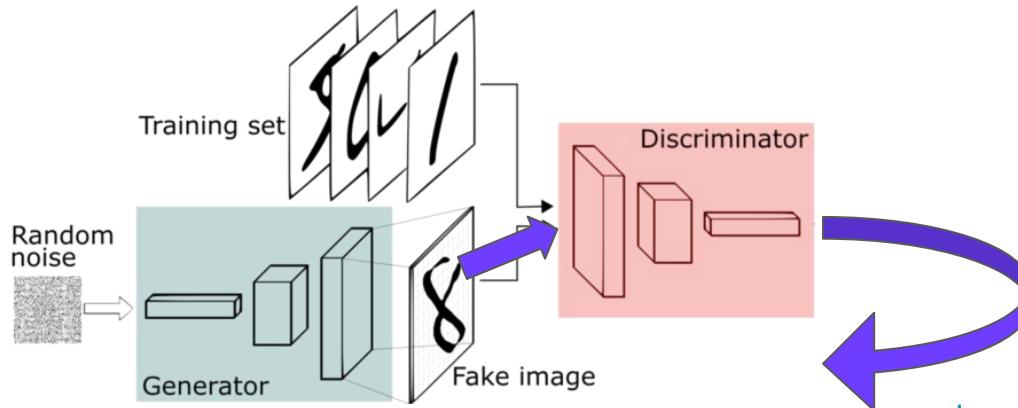
# GANs: Generative Adversarial Networks

## Training the Discriminator

- Straightforward binary classification (Real vs. Fake)

## Training the Generator

- Generator produces an image → pass it through the Discriminator
- If Discriminator classifies it as Fake → update the Generator to produce better images



[Image source](#)

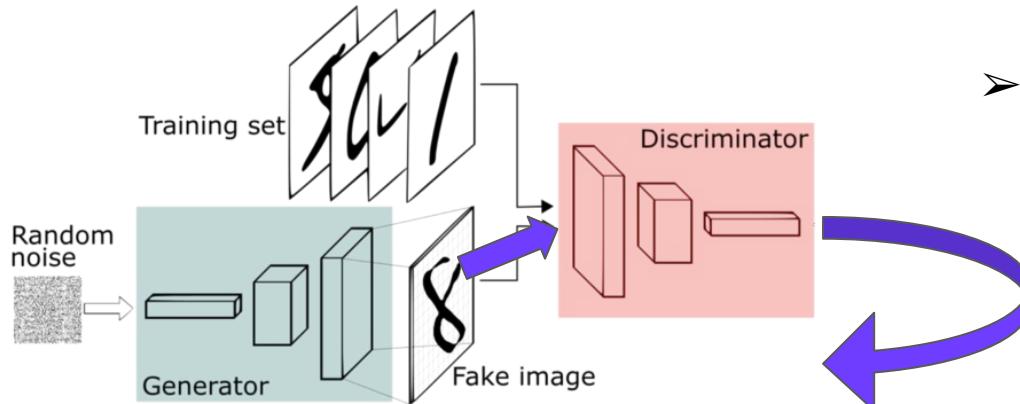
# GANs: Generative Adversarial Networks

## Training the Discriminator

- Straightforward binary classification (Real vs. Fake)

## Training the Generator

- Generator produces an image → pass it through the Discriminator
- If Discriminator classifies it as Fake → update the Generator to produce better images



## What's the objective?

- Generator gets good enough that Discriminator can't tell the difference between Real and Fake

Image source

# GANs: Generative Adversarial Networks

- Real data:  $x \sim p(x)$
- Generated data:  $x \sim p_G(x)$  (Actually,  $x \sim p_G(x|z)$ )
- Goal:  $p_G(x) \approx p(x)$
- Discriminator D is trained to assign the correct label (Real vs Fake)
- Generator G is trained to produce images that will fool the Discriminator D

# GANs: Generative Adversarial Networks

- Real data:  $x \sim p(x)$
- Generated data:  $x \sim p_G(x)$  (Actually,  $x \sim p_G(x|z)$ )
- Goal:  $p_G(x) \approx p(x)$
- Discriminator D is trained to assign the correct label (Real vs Fake)
- Generator G is trained to produce images that will fool the Discriminator D

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

# GANs: Generative Adversarial Networks

- Real data:  $x \sim p(x)$
- Generated data:  $x \sim p_G(x)$  (Actually,  $x \sim p_G(x|z)$ )
- Goal:  $p_G(x) \approx p(x)$
- Discriminator D is trained to assign the correct label (Real vs Fake)
- Generator G is trained to produce images that will fool the Discriminator D

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Binary cross entropy:  $-(y \log(p) + (1 - y) \log(1 - p))$

# How good can GANs get?

# How good can GANs get?

Really good. So good, it's freaky!

Example: <https://thispersondoesnotexist.com/>

# How good can GANs get?

Really good. So good, it's freaky!

Example: <https://thispersondoesnotexist.com/>



# How good can GANs get?

Really good. So good, it's freaky!

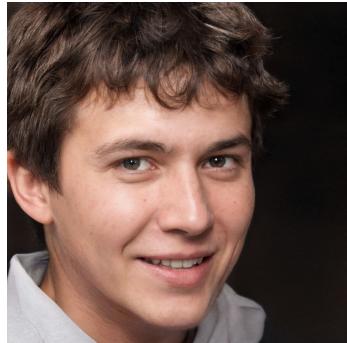
Example: <https://thispersondoesnotexist.com/>



# How good can GANs get?

Really good. So good, it's freaky!

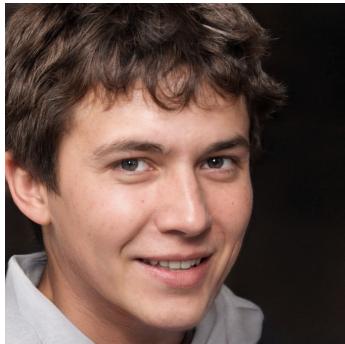
Example: <https://thispersondoesnotexist.com/>



# How good can GANs get?

Really good. So good, it's freaky!

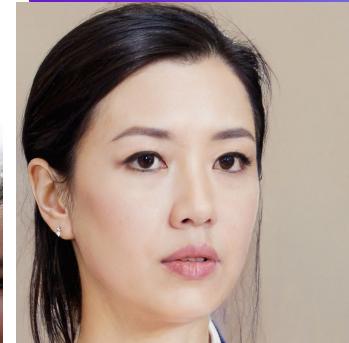
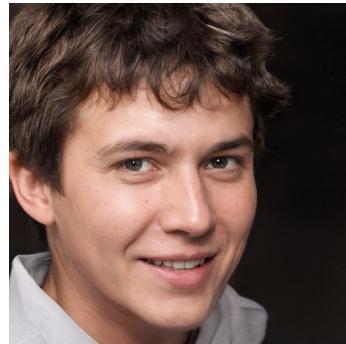
Example: <https://thispersondoesnotexist.com/>



# How good can GANs get?

Really good. So good, it's freaky!

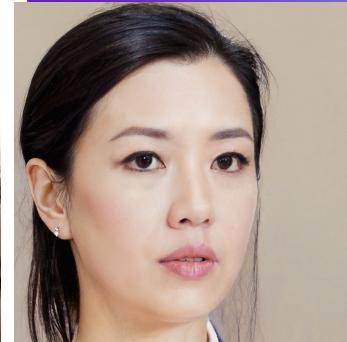
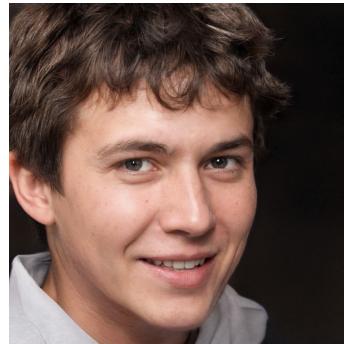
Example: <https://thispersondoesnotexist.com/>



# How good can GANs get?

Really good. So good, it's freaky!

Example: <https://thispersondoesnotexist.com/>



Good for these unsupervised GANs. But, how can they help us with semi-supervised learning?

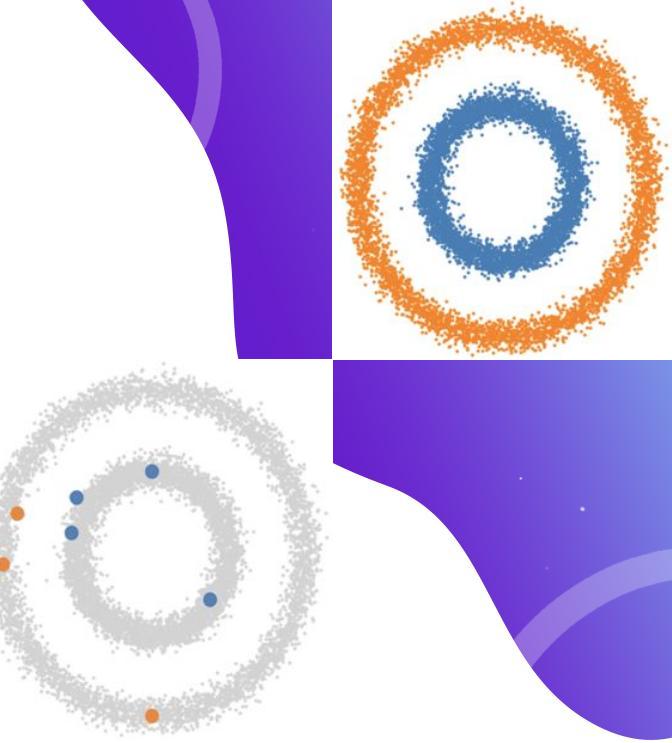
# THEORY

## Semi-Supervised Learning with GANs

# WHAT IS THE BASIC IDEA HERE?

Ref: [Improved Techniques for Training GANs, 2016](#)

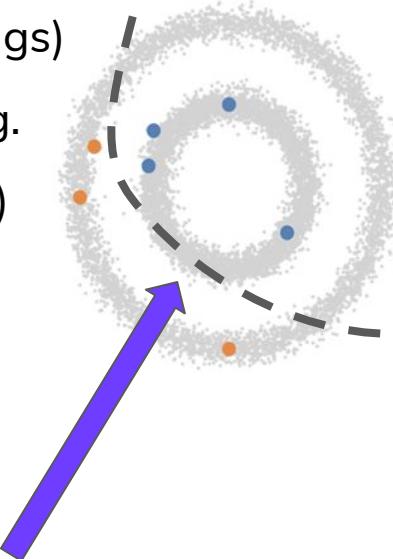
- Our data has a certain distribution (e.g. the two rings)
- The GAN will attempt to learn that distribution (e.g. Generator will be trained to produce the two rings)
- The few labeled datapoints can then be used to assign class labels (e.g. **blue** ring vs. **orange** ring)



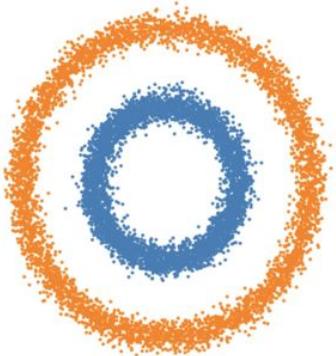
# WHAT IS THE BASIC IDEA HERE?

Ref: [Improved Techniques for Training GANs, 2016](#)

- Our data has a certain distribution (e.g. the two rings)
- The GAN will attempt to learn that distribution (e.g. Generator will be trained to produce the two rings)
- The few labeled datapoints can then be used to assign class labels (e.g. **blue** ring vs. **orange** ring)



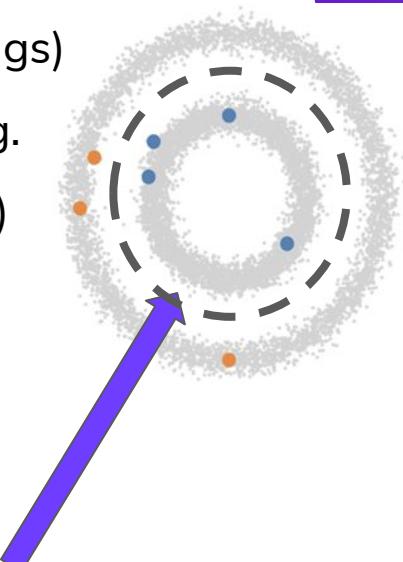
Decision boundary **without** the unlabelled data



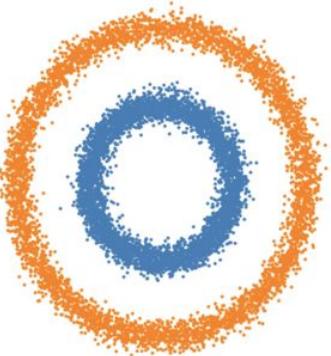
# WHAT IS THE BASIC IDEA HERE?

Ref: [Improved Techniques for Training GANs, 2016](#)

- Our data has a certain distribution (e.g. the two rings)
- The GAN will attempt to learn that distribution (e.g. Generator will be trained to produce the two rings)
- The few labeled datapoints can then be used to assign class labels (e.g. **blue** ring vs. **orange** ring)



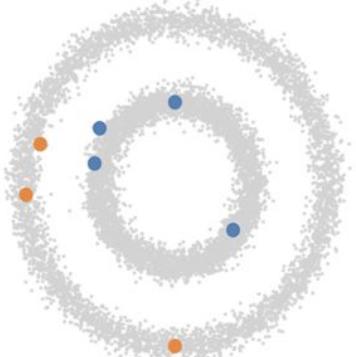
Decision boundary **with** the unlabelled data



# HOW DOES THIS WORK?

Ref: [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

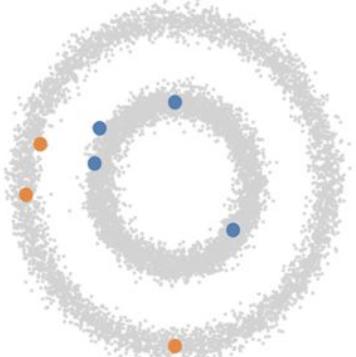
- Unsupervised GAN: Generator + Discriminator (Real vs. Fake)
- Supervised classifier: K classes (blues and oranges)
- Semi-supervised: Discriminator = Classifier into  $K+1$  classes
- The first  $K$  classes are Real; the added ( $K+1$ ) class is Fake  
(blues, oranges and Fakes)
- Semi-supervised GAN:  
Generator + the new Discriminator/Classifier



# HOW DOES THIS WORK?

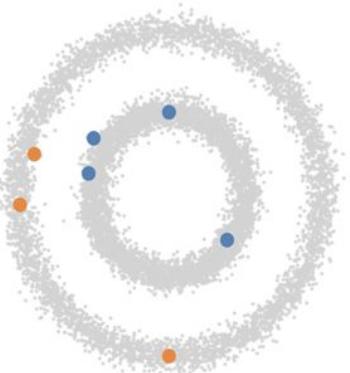
Ref: [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

- Unsupervised GAN: Generator + Discriminator (Real vs. Fake)
- Supervised classifier: K classes (blues and oranges)
- Semi-supervised: Discriminator = Classifier into  $K+1$  classes
- The first  $K$  classes are Real; the added ( $K+1$ ) class is Fake  
(blues, oranges and Fakes)
- Semi-supervised GAN:  
Generator + the new Discriminator/Classifier



# HOW DOES THIS WORK?

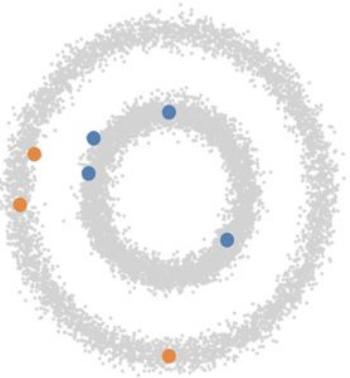
Ref: [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)



- Unsupervised GAN: Generator + Discriminator (Real vs. Fake)
- Supervised classifier: K classes (**blues** and **oranges**)
- Semi-supervised: Discriminator = Classifier into  $K+1$  classes
- The first  $K$  classes are Real; the added ( $K+1$ ) class is Fake  
(blues, oranges and Fakes)
- Semi-supervised GAN:  
Generator + the new Discriminator/Classifier

# HOW DOES THIS WORK?

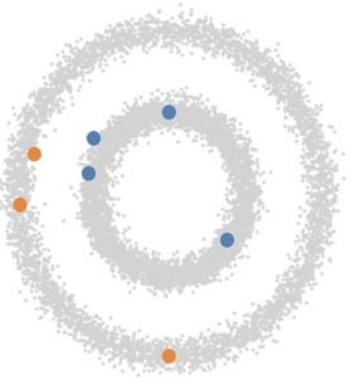
Ref: [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)



- Unsupervised GAN: Generator + Discriminator (Real vs. Fake)
- Supervised classifier: K classes (**blues** and **oranges**)
- Semi-supervised: Discriminator = Classifier into K+1 classes
- The first K classes are Real; the added (K+1) class is Fake  
(blues, oranges and Fakes)
- Semi-supervised GAN:  
Generator + the new Discriminator/Classifier

# HOW DOES THIS WORK?

Ref: [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

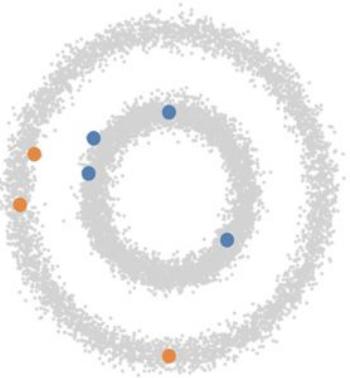


- Unsupervised GAN: Generator + Discriminator (Real vs. Fake)
- Supervised classifier: K classes (**blues** and **oranges**)
- Semi-supervised: Discriminator = Classifier into K+1 classes
- The first K classes are Real; the added (K+1) class is Fake  
(**blues**, **oranges** and Fakes)
- Semi-supervised GAN:

Generator + the new Discriminator/Classifier

# HOW DOES THIS WORK?

Ref: [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)



- Unsupervised GAN: Generator + Discriminator (Real vs. Fake)
- Supervised classifier: K classes (**blues** and **oranges**)
- Semi-supervised: Discriminator = Classifier into K+1 classes
- The first K classes are Real; the added (K+1) class is Fake  
(**blues**, **oranges** and Fakes)
- Semi-supervised GAN:  
**Generator + the new Discriminator/Classifier**

# HOW ARE THE NETWORKS TRAINED?

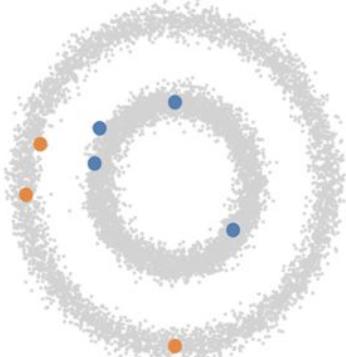
Three kinds of data: labeled, unlabeled, and generated from noise

Discriminator is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (**blues** or **oranges**, does not matter which)
- generated data to (K+1) class (**Fake**)

Generator is trained to:

- generate images that Discriminator would assign to one of K classes (**blues** or **oranges**, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

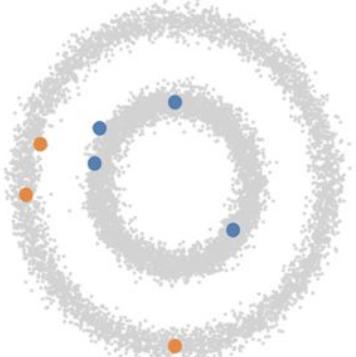
Three kinds of data: labeled, unlabeled, and generated from noise

Discriminator is trained to assign:

- labeled data to correct classes (**blues vs. oranges**)
- unlabeled data to one of K classes (**blues or oranges**, does not matter which)
- generated data to (K+1) class (**Fake**)

Generator is trained to:

- generate images that Discriminator would assign to one of K classes (**blues or oranges**, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

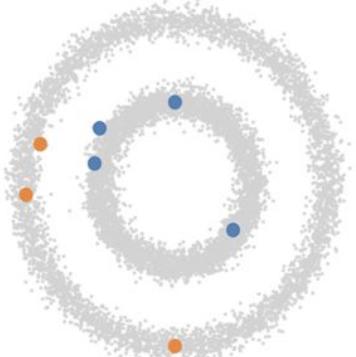
Three kinds of data: **labeled**, unlabeled, and generated from noise

Discriminator is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (**blues** or **oranges**, does not matter which)
- generated data to (K+1) class (**Fake**)

Generator is trained to:

- generate images that Discriminator would assign to one of K classes (**blues** or **oranges**, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

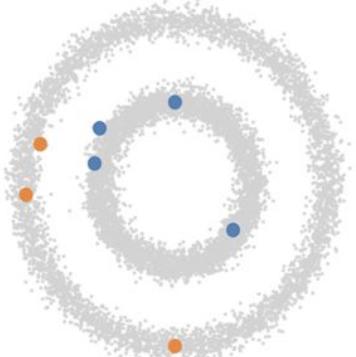
Three kinds of data: **labeled**, **unlabeled**, and generated from noise

Discriminator is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (**blues** or **oranges**, does not matter which)
- generated data to (K+1) class (**Fake**)

Generator is trained to:

- generate images that Discriminator would assign to one of K classes (**blues** or **oranges**, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

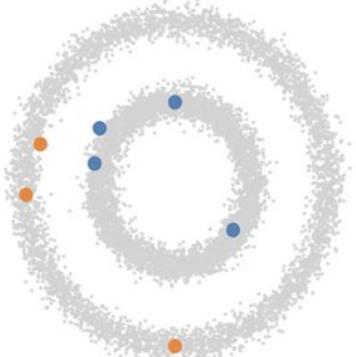
Three kinds of data: labeled, unlabeled, and **generated** from noise

Discriminator is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (**blues** or **oranges**, does not matter which)
- generated data to (K+1) class (**Fake**)

Generator is trained to:

- generate images that Discriminator would assign to one of K classes (**blues** or **oranges**, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

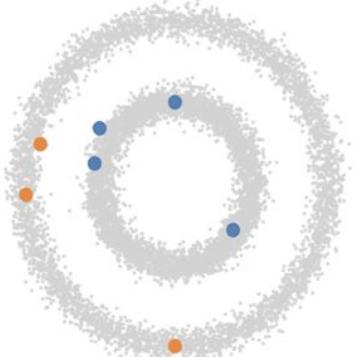
Three kinds of data: labeled, unlabeled, and generated from noise

**Discriminator** is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (blues or oranges, does not matter which)
- generated data to (K+1) class (Fake)

**Generator** is trained to:

- generate images that Discriminator would assign to one of K classes (blues or oranges, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

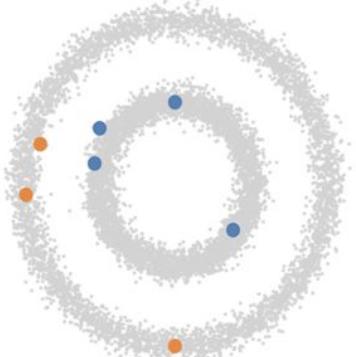
Three kinds of data: labeled, unlabeled, and generated from noise

**Discriminator** is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (**blues** or **oranges**, does not matter which)
- generated data to (K+1) class (Fake)

**Generator** is trained to:

- generate images that Discriminator would assign to one of K classes (**blues** or **oranges**, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

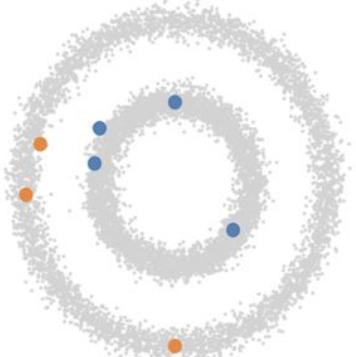
Three kinds of data: labeled, unlabeled, and generated from noise

**Discriminator** is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (**blues** or **oranges**, does not matter which)
- generated data to (K+1) class (Fake)

**Generator** is trained to:

- generate images that Discriminator would assign to one of K classes (**blues** or **oranges**, does not matter which)



# HOW ARE THE NETWORKS TRAINED?

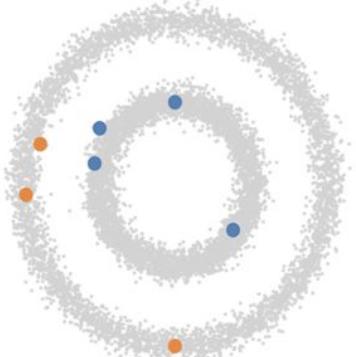
Three kinds of data: labeled, unlabeled, and generated from noise

**Discriminator** is trained to assign:

- labeled data to correct classes (**blues** vs. **oranges**)
- unlabeled data to one of K classes (**blues** or **oranges**, does not matter which)
- generated data to (K+1) class (Fake)

**Generator** is trained to:

- generate images that Discriminator would assign to one of K classes (**blues** or **oranges**, does not matter which)



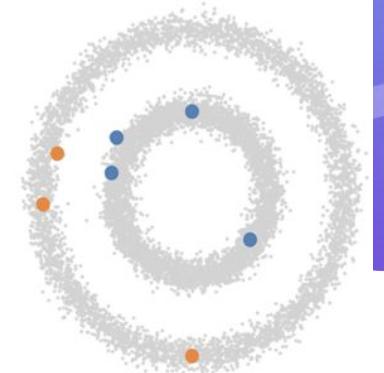
# HOW GOOD IS THE TRAINED GENERATOR?

Ref: [Good Semi-supervised Learning that Requires a Bad GAN, 2017](#)

- Unsupervised GANs can produce very impressive images
- However, the GANs used for semi-supervised learning do not
- To get a good classifier, the Generator should be “bad” = **complement** (generates complement samples in feature space)

The two rings example:

- Generating more points for the two circles would not help us
- Generating Fake points between circles  
→ better decision boundary between classes



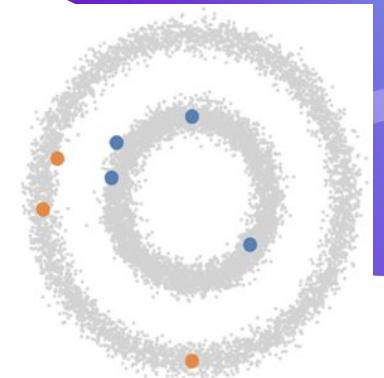
# HOW GOOD IS THE TRAINED GENERATOR?

Ref: [Good Semi-supervised Learning that Requires a Bad GAN, 2017](#)

- Unsupervised GANs can produce very impressive images
- However, the GANs used for semi-supervised learning **do not**
- To get a good classifier, the Generator should be “bad” = **complement** (generates complement samples in feature space)

The two rings example:

- Generating more points for the two circles would not help us
- Generating **Fake** points between circles  
→ better decision boundary between classes



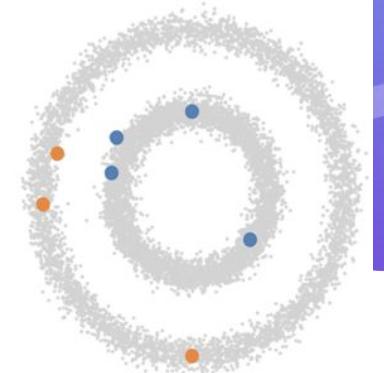
# HOW GOOD IS THE TRAINED GENERATOR?

Ref: [Good Semi-supervised Learning that Requires a Bad GAN, 2017](#)

- Unsupervised GANs can produce very impressive images
- However, the GANs used for semi-supervised learning **do not**
- To get a good classifier, the Generator should be “bad” = **complement** (generates complement samples in feature space)

The two rings example:

- Generating more points for the two circles would not help us
- Generating **Fake** points between circles  
→ better decision boundary between classes



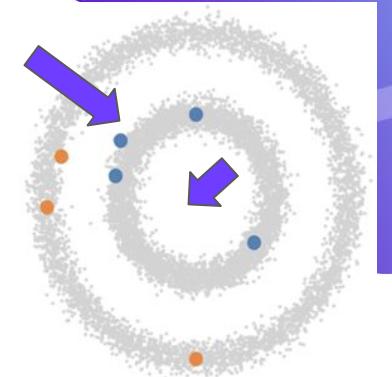
# HOW GOOD IS THE TRAINED GENERATOR?

Ref: [Good Semi-supervised Learning that Requires a Bad GAN, 2017](#)

- Unsupervised GANs can produce very impressive images
- However, the GANs used for semi-supervised learning **do not**
- To get a good classifier, the Generator should be “bad” = **complement** (generates complement samples in feature space)

The two rings example:

- Generating more points for the two circles would not help us
- Generating **Fake** points between circles  
→ better decision boundary between classes



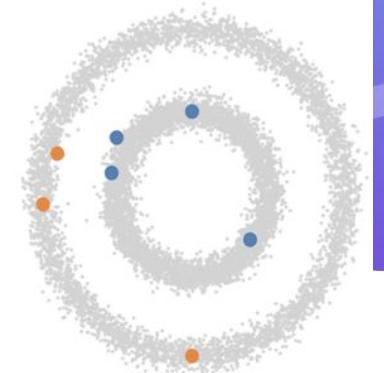
# HOW GOOD IS THE TRAINED GENERATOR?

Ref: [Good Semi-supervised Learning that Requires a Bad GAN, 2017](#)

- Unsupervised GANs can produce very impressive images
- However, the GANs used for semi-supervised learning **do not**
- To get a good classifier, the Generator should be “bad” = **complement** (generates complement samples in feature space)

The two rings example:

- Generating more points for the two circles would not help us
- Generating *Fake* points between circles  
→ better decision boundary between classes



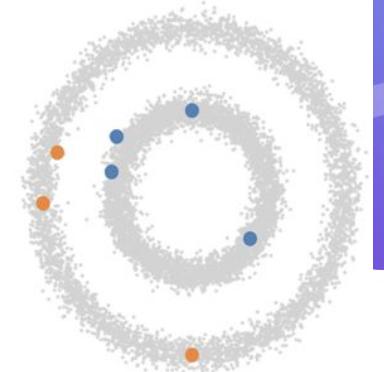
# HOW GOOD IS THE TRAINED GENERATOR?

Ref: [Good Semi-supervised Learning that Requires a Bad GAN, 2017](#)

- Unsupervised GANs can produce very impressive images
- However, the GANs used for semi-supervised learning **do not**
- To get a good classifier, the Generator should be “bad” = **complement** (generates complement samples in feature space)

The two rings example:

- Generating more points for the two circles would not help us
- Generating Fake points between circles  
→ better decision boundary between classes



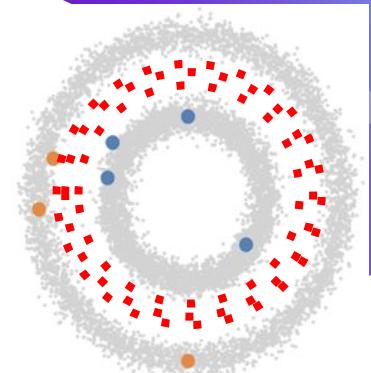
# HOW GOOD IS THE TRAINED GENERATOR?

Ref: [Good Semi-supervised Learning that Requires a Bad GAN, 2017](#)

- Unsupervised GANs can produce very impressive images
- However, the GANs used for semi-supervised learning **do not**
- To get a good classifier, the Generator should be “bad” = **complement** (generates complement samples in feature space)

The two rings example:

- Generating more points for the two circles would not help us
- Generating Fake points between circles  
→ better decision boundary between classes



# PRACTICE

## A TALE OF CATS & DOGS

# PRACTICE

## A TALE OF CATS & DOGS

[github.com/opetrova/SemiSupervisedPytorchGAN](https://github.com/opetrova/SemiSupervisedPytorchGAN)

# PRACTICE

A TALE OF CATS & DOGS

[github.com/opetrova/SemiSupervisedPytorchGAN](https://github.com/opetrova/SemiSupervisedPytorchGAN)

CAUTIONARY



# MORAL(S) OF THE STORY

# MORAL(S) OF THE STORY

1. Your labelled data matters **A LOT**



10% average rise in accuracy for randomly chosen cats+dogs  
vs. 20%+ for a curated labeled dataset (“representative” images)

2. Naively adding a new Fake class to the Classifier does **NOT** work

Instead:

- Keep number of output logits at K (K=2 for binary classification)
- Change your loss function:

Ref: [Improved Techniques for Training GANs, 2016](#), [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

- The idea: output a strong class prediction for Real examples, and small class predictions for Fake examples

# MORAL(S) OF THE STORY

1. Your labelled data matters **A LOT**



10% average rise in accuracy for randomly chosen cats+dogs  
vs. 20%+ for a curated labeled dataset (“representative” images)

2. Naively adding a new Fake class to the Classifier does **NOT** work

Instead:

- Keep number of output logits at K (K=2 for binary classification)
- Change your loss function:



Ref: [Improved Techniques for Training GANs, 2016](#), [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

- The idea: output a strong class prediction for Real examples, and small class predictions for Fake examples

# MORAL(S) OF THE STORY

1. Your labelled data matters **A LOT**



10% average rise in accuracy for randomly chosen cats+dogs  
vs. 20%+ for a curated labeled dataset (“representative” images)

2. Naively adding a new Fake class to the Classifier does **NOT** work



**Instead:**

- Keep number of output logits at K (K=2 for binary classification)
- Change your loss function:

Ref: [Improved Techniques for Training GANs, 2016](#), [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

- The idea: output a strong class prediction for Real examples, and small class predictions for Fake examples

# MORAL(S) OF THE STORY

1. Your labelled data matters **A LOT**



10% average rise in accuracy for randomly chosen cats+dogs  
vs. 20%+ for a curated labeled dataset (“representative” images)

2. Naively adding a new Fake class to the Classifier does **NOT** work



**Instead:**

- Keep number of output logits at K (K=2 for binary classification)
- Change your loss function:  $D(\mathbf{x}) = \frac{Z(\mathbf{x})}{Z(\mathbf{x})+1}$ , where  $Z(\mathbf{x}) = \sum_{k=1}^K \exp[l_k(\mathbf{x})]$

Ref: [Improved Techniques for Training GANs, 2016](#), [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

- The idea: output a strong class prediction for Real examples, and small class predictions for Fake examples

# MORAL(S) OF THE STORY

1. Your labelled data matters **A LOT**



10% average rise in accuracy for randomly chosen cats+dogs  
vs. 20%+ for a curated labeled dataset (“representative” images)

2. Naively adding a new Fake class to the Classifier does **NOT** work



**Instead:**

- Keep number of output logits at K (K=2 for binary classification)
- Change your loss function:  $D(\mathbf{x}) = \frac{Z(\mathbf{x})}{Z(\mathbf{x})+1}$ , where  $Z(\mathbf{x}) = \sum_{k=1}^K \exp[l_k(\mathbf{x})]$

Ref: [Improved Techniques for Training GANs, 2016](#), [Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks, 2016](#)

- The idea: output a strong class prediction for Real examples, and small class predictions for Fake examples

# MERCI

**CONTACT**  
Olga PETROVA

<https://www.linkedin.com/in/olga-p-petrova/>